

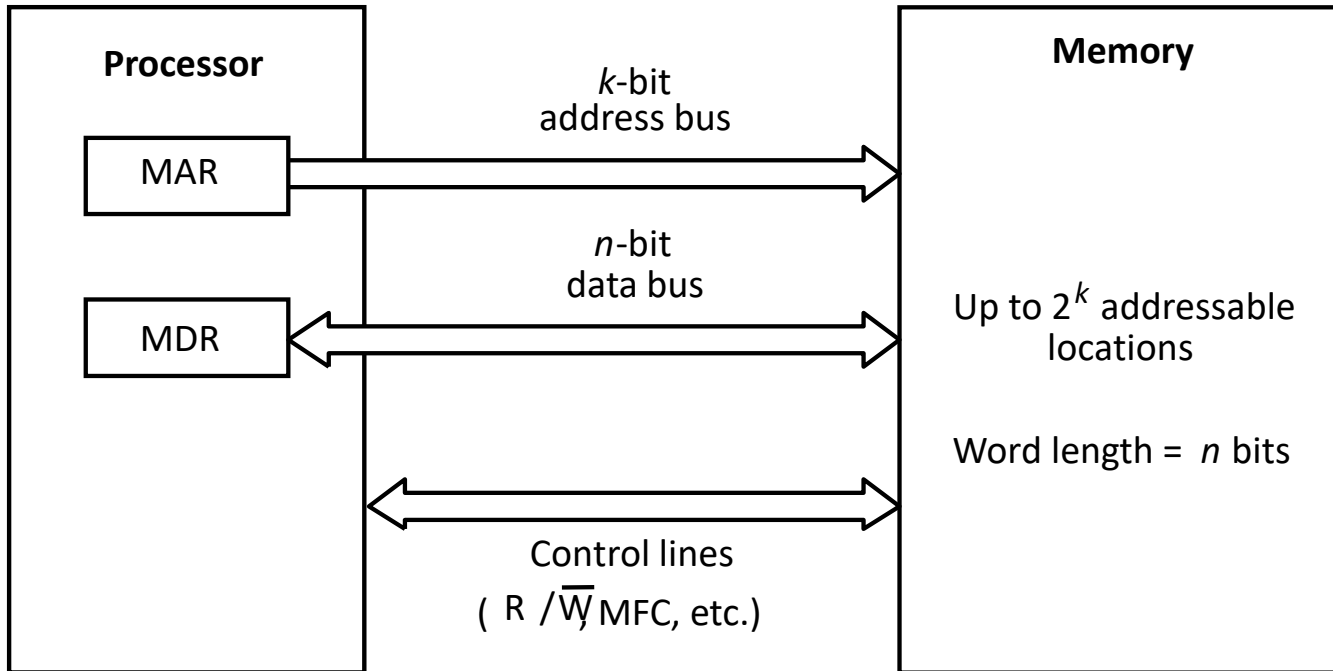
# The Memory System

(Chapter 5: Hamacher, Zaky)

## Basic concepts [SELF STUDY]

- Maximum size of the memory depends on the addressing scheme:
  - ◆ 16-bit computer generates 16-bit addresses and can address up to  $2^{16}$  memory locations.
  - ◆ Number of locations represents the size of the address space of a computer.
- Most modern computers are byte-addressable.
- Memory is designed to store and retrieve data in word-length quantities.
  - ◆ Word length of a computer is commonly defined as the number of bits stored and retrieved in one memory operation.

## Basic concepts (contd..) [SELF STUDY]



Recall that the **data** transfers between a processor and memory involves two registers **MAR** and **MDR**.

If the **address bus** is  $k$ -bits, then the length of **MAR** is  $k$  bits.

If the **word length** is  $n$ -bits, then the length of **MDR** is  $n$  bits.

**Control** lines include  **$R/\overline{W}$**  and **MFC**.

For Read operation  $R/\overline{W} = 1$  and for Write operation  $R/\overline{W} = 0$ .

## Basic concepts (contd..)

### □ Measures for the **speed** of a **memory**:

- ◆ Elapsed time between the initiation of an operation and the completion of an operation is the memory access time (e.g. the time between the **Read** & **MFC** signals).
- ◆ Minimum time between the initiation of two successive memory operations is memory cycle time (e.g. the time between two successive **Read** operations)
- ◆ Memory Cycle time is slightly longer than memory access time

### □ In general, the **faster** a memory system, the **costlier** it is and the **smaller** it is.

### □ An important design issue is to **provide** a computer system with as **large** and **fast** a memory as possible, within a given **cost** target.

### □ Several **techniques** to **increase** the effective **size** and **speed** of the memory:

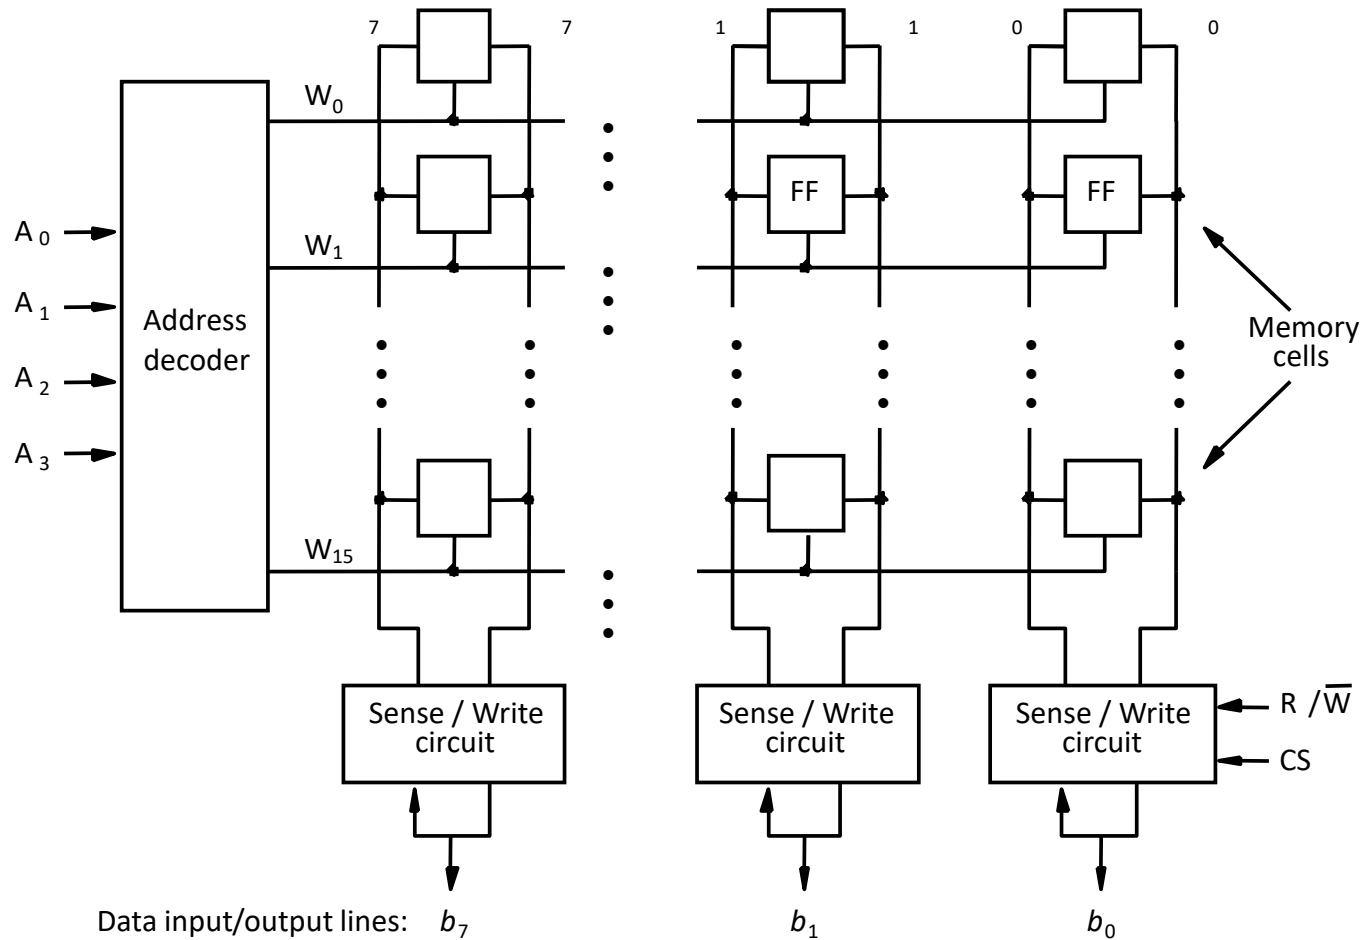
- ◆ **Cache** memory (to increase the effective **speed**).
- ◆ **Virtual** memory (to increase the effective **size**).

# Semiconductor RAM memories

- ❑ Random Access Memory (RAM) memory unit is a unit where any location can be addressed in a fixed amount of time, independent of the location's address.
- ❑ Internal organization of memory chips:
  - ◆ Each memory cell can hold one bit of information.
  - ◆ Memory cells are organized in the form of an array.
  - ◆ One row is one memory word.
  - ◆ All cells of a row are connected to a common line, known as the "word line".
  - ◆ Word line is connected to the address decoder.
  - ◆ Sense/write circuits are connected to the data input/output lines of the memory chip.

# Semiconductor RAM memories (contd..)

Internal organization of memory chips size=128 bits (16x8)



# Semiconductor RAM memories (contd..)

## □ Static RAMs (SRAMs):

- ◆ Consist of **circuits** that are capable of **retaining** their **state** as long as the power is applied.
- ◆ **Volatile memories**, because their contents are lost when power is interrupted.
- ◆ **Access times** of static RAMs are in the range of few **nanoseconds**.
- ◆ **Capacity** is **low** (each **cell** consists of **6 transistors**)
- ◆ However, the **cost** is usually **high**.

## □ Dynamic RAMs (DRAMs):

- ◆ Do **not retain** their **state** indefinitely.
- ◆ **Contents** must be periodically **refreshed**.
- ◆ Contents may be refreshed while accessing them for reading.
- ◆ **Access time** is **longer** than **SRAM**
- ◆ **Capacity** is **higher** than **SRAM** (each **cell** consists of **1 transistor**)
- ◆ **Cost** is **lower** than **SRAM**

## Semiconductor RAM memories (contd..)

- ❑ Data is transferred between the processor and the memory in units of single word or a block.
- ❑ Speed and efficiency of data transfer has a large impact on the performance of a computer.
- ❑ Two metrics of performance: Latency and Bandwidth.
- ❑ Latency:
  - ◆ Time taken to transfer a single word of data to or from memory.
  - ◆ Definition is clear if the memory operation involves transfer of a single word of data.
  - ◆ In case of a block transfer, latency is the time it takes to transfer first word of data.
  - ◆ Time required to transfer first word in a block is substantially larger than the time required to transfer consecutive words in a block.



## Semiconductor RAM memories (contd..)

- ❑ How much time is needed to transfer a single block of data.
- ❑ Blocks can be variable in size, it is useful to define a performance measure in terms of the number of bits or bytes transferred in one second.
- ❑ This performance measure is referred to as memory bandwidth.
- ❑ Bandwidth of a memory unit depends on:
  - ◆ Speed of access to the chip.
  - ◆ How many bits can be accessed in parallel.
- ❑ Bandwidth of data transfer between processor and memory unit also depends on:
  - ◆ Transfer capability of the links that connect the processor and memory, or the speed of the bus.

# Memory systems

- ❑ Various factors such as **cost**, **speed**, **power** consumption and **size** of the chip determine how a **RAM** is chosen for a given application.
- ❑ **Static RAMs:**
  - ◆ Chosen when **speed** is the primary concern.
  - ◆ **Circuit** implementing the basic cell is highly **complex**, so **cost** and **size** are affected.
  - ◆ Used mostly in **cache** memories.
- ❑ **Dynamic RAMs:**
  - ◆ Predominantly used for **implementing computer main memories**.
  - ◆ **High densities** available in these chips.
  - ◆ **Economically viable** for implementing large memories.

# Read-Only Memories (ROMs)

- ❑ SRAM and SDRAM chips are **volatile**:
  - ◆ Lose the contents when the power is turned off.
- ❑ Many applications need memory devices to retain contents after the power is turned off.
  - ◆ For **example**, computer is turned on, the **operating system** must be **loaded** from the **disk** into the **memory**.
  - ◆ **Store instructions** which would **load** the **OS** from the **disk**.
  - ◆ Need to store these instructions so that they will not be lost after the power is turned off.
  - ◆ We need to **store** the **instructions** into a **non-volatile memory**.
- ❑ Non-volatile memory is read in the same manner as volatile memory.
  - ◆ **Separate writing process** is needed to place information in this memory.
  - ◆ Normal operation involves only reading of data, this type of memory is called **Read-Only memory (ROM)**.

# Read-Only Memories (contd..) [SELF STUDY]

## □ Read-Only Memory:

- ◆ Data are written into a ROM when it is manufactured.

## □ Programmable Read-Only Memory (PROM):

- ◆ Allow the data to be loaded by a user.
- ◆ Process of inserting the data is irreversible.
- ◆ Storing information specific to a user in a ROM is expensive.
- ◆ Providing programming capability to a user may be better.

## □ Erasable Programmable Read-Only Memory (EPROM):

- ◆ Stored data to be erased and new data to be loaded.
- ◆ Flexibility, useful during the development phase of digital systems.
- ◆ Erasable, reprogrammable ROM.
- ◆ Erasure requires exposing the ROM to UV light.

## Read-Only Memories (contd..) [SELF STUDY]

### □ Electrically Erasable Programmable Read-Only Memory (EEPROM):

- ◆ To **erase** the contents of **EPROMs**, they have to be **exposed** to **ultraviolet** light.
- ◆ **Physically removed** from the circuit.
- ◆ **EEPROMs** the contents can be **stored** and **erased electrically**.

# Flash memory [SELF STUDY]

- ❑ Flash memory has similar approach to EEPROM.
- ❑ Read the contents of a single cell, but write the contents of an entire block of cells.
- ❑ Flash devices have greater density.
  - ◆ Higher capacity and low storage cost per bit.
- ❑ Power consumption of flash memory is very low, making it attractive for use in equipment that is battery-driven.
- ❑ Single flash chips are not sufficiently large, so larger memory modules are implemented using flash cards and flash drives.

# Acronyms [SELF STUDY]

RAM	--Random Access Memory	time taken to access any arbitrary location in memory is constant (c.f., disks)
ROM	--Read Only Memory	ROMs are RAMs which can only be written to once; thereafter they can only be read Older uses included storage of bootstrap info
PROM	--Programmable ROM	A ROM which can be bench programmed
EPROM	--Erasable PROM	A PROM which can be erased for rewriting
EEPROM	--Electrically EPROM	A PROM which can be erased electrically.
SRAM	--Static RAM	RAM chip which loses contents upon power off
DRAM	--Dynamic RAM	RAM chip whose contents need to be refreshed.
SDRAM	--Synchronous DRAM	RAM whose memory operations are synchronized with a clock signal.

## Memory hierarchy

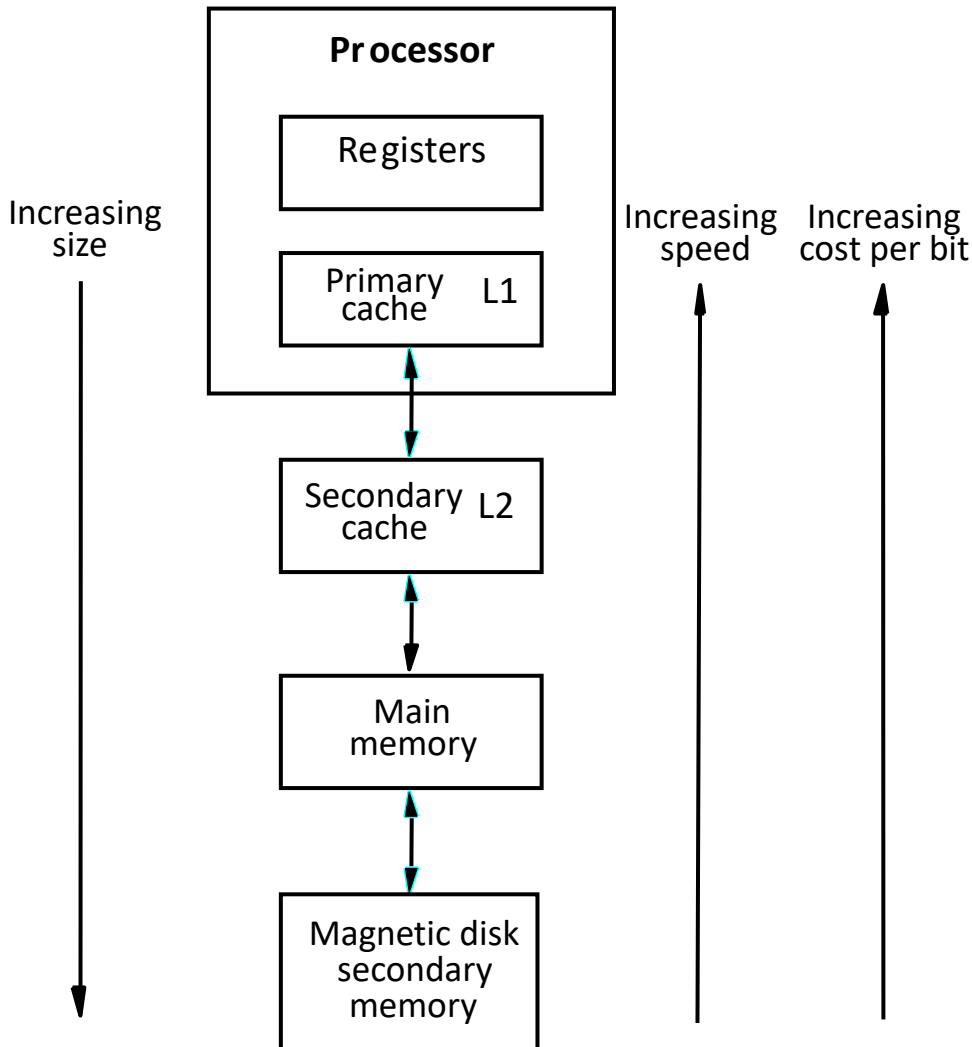
- A big challenge in the design of a computer system is to provide a sufficiently large memory, with a reasonable speed at an affordable cost.
- Static RAM:
  - Very **fast**, but **expensive**, because a basic SRAM cell has a **complex circuit** making it impossible to pack a large number of cells onto a single chip.
- Dynamic RAM:
  - **Simpler** basic cell **circuit**, hence are much **less expensive**, but significantly **slower** than SRAMs.
- Magnetic disks:
  - Storage provided by DRAMs is higher than SRAMs, but is still less than what is necessary.
  - Secondary storage such as magnetic disks provide a **large** amount of **storage**, but is much **slower** than DRAMs.



# Memory hierarchy (contd..)

- All these types of memory units are employed effectively in a computer.
  - SRAM -- Smaller units where speed is critical, namely, cache memories.
  - DRAM -- Large, yet affordable memory, namely, main memory.
  - Magnetic disks -- Huge amount of cost-effective storage.
- Computer memory can be viewed as a hierarchy.

# Memory hierarchy (contd..)



- **Fastest** access is to the data held in processor **registers**. Registers are at the top of the memory hierarchy.
- Relatively small amount of memory that can be implemented on the processor chip. This is **processor cache**.
- Two levels of cache:
  - Level 1 (**L1**) cache is on the processor chip.
  - Level 2 (**L2**) cache is in between main memory and processor.
- Next level is **main memory**, implemented as SIMMs. Much larger, but much slower than cache memory.
- Next level is **magnetic disks**. Huge amount of inexpensive storage.
- **Speed of memory access is critical, the idea is to bring instructions and data that will be used in the near future as close to the processor as possible.**

## Cache memories

- Processor is much faster than the main memory.
  - As a result, the **processor** has to spend much of its time **waiting** while **instructions** and **data** are being **fetch**ed from the main **memory**.
  - Major obstacle towards achieving good performance.
- Speed of the main memory cannot be increased beyond a certain point.
- Cache memory is an architectural arrangement which makes the main memory appear faster to the processor than it really is.
- Cache memory is based on the property of computer programs known as “locality of reference”.

## Locality of reference

- Analysis of programs indicates that many instructions in localized areas of a program are executed repeatedly during some period of time, while the others are accessed relatively less frequently.
  - These instructions may be the ones in a loop, nested loop or few procedures calling each other repeatedly.
  - This is called “locality of reference”.
- **Temporal locality of reference:**
  - Recently executed instruction is likely to be executed again very soon.
- **Spatial locality of reference:**
  - Instructions with addresses close to a recently instruction are likely to be executed soon.


# Locality of reference (contd..)

- Cache memory is based on the concept of locality of reference.
  - If **active segments** of a **program** are **placed** in a fast **cache** memory, then the **execution time** can be **reduced**.
- **Temporal** locality of reference:
  - Whenever an **instruction** or **data** is needed for the **first time**, it should be brought into a **cache**. It will hopefully be **used again** repeatedly.
- **Spatial** locality of reference:
  - Instead of fetching just one item from the main memory to the cache at a time, **several items** that have **addresses adjacent** to the **item** being **fetches** may be **useful**.
  - The term “**block**” refers to a **set of contiguous addresses** locations of some size.

# Hit and Miss [SELF STUDY]

- Focus on *any two adjacent* levels – called, *upper* (closer to CPU) and *lower* (farther from CPU) – in the memory hierarchy, because each block copy is always between two adjacent levels
- **Terminology:**
  - *block*: minimum unit of data to move between levels (Usually 1 word)
  - *hit*: data requested is in upper level (example: data found in cache!)
  - *miss*: data requested is not in upper level (e.g., data not in cache, So, go to lower level (main memory or hard disk) to find that data)
  - *hit rate*: fraction of memory accesses that are hits (i.e., found at upper level, say cache)
  - *miss rate*: fraction of memory accesses that are not hits
- $\text{miss rate} = 1 - \text{hit rate}$
- *hit time*: time to determine if the access is really a hit + time to get and deliver the data from the upper level (e.g., cache) to the CPU
- *miss penalty*: time to determine if the access is a miss + time to replace block at upper level (say, cache) with corresponding block at lower level (e.g., RAM) + time to deliver the block to the CPU

# Cache memories



- *Processor issues a Read request, a block of words is transferred from the main memory to the cache, one word at a time.*
- *Subsequent references to the data in this block of words are found in the cache.*
- *At any given time, only some blocks in the main memory are held in the cache.*
- *Which blocks in the main memory are in the cache is determined by a “mapping function”.*
- *When the cache is full, and a block of words needs to be transferred from the main memory, some block of words in the cache must be replaced. This is determined by a “replacement algorithm”.*

# Cache memories (contd..)

- Existence of a *cache* is *transparent* to the *processor*. The processor issues *Read* and *Write* requests in the *same* manner.
- If the *data* is in the *cache* it is called a Read or Write **hit**.
- **Read hit:**
  - The *data* is obtained from the *cache*.
- **Write hit:**
  - *Cache* is a *replica* of the *contents* of the *main memory*.
  - *Contents* of the *cache* and the *main memory* may be *updated simultaneously*.  
This is the write-through protocol.
  - *Update* the *contents* of the *cache*, and *mark* it as *updated* by setting a *bit* known as the dirty bit or modified bit.  
The *contents* of the *main memory* are *updated* when this *block* is *replaced*.  
This is write-back or copy-back protocol.



# Cache memories (contd..)

- If the *data is not present in the cache*, then a Read miss or Write miss occurs.

- **Read miss:**

- *Block of words containing this requested word is transferred from the memory.*
- *After the block is transferred, the desired word is forwarded to the processor.*
- *The desired word may also be forwarded to the processor as soon as it is transferred without waiting for the entire block to be transferred. This is called load-through or early-restart.*

- **Write-miss:**


- **Write-through** protocol is used, then the contents of the main memory are updated directly.
- If **write-back** protocol is used, the block containing the addressed word is first brought into the cache. The desired word is overwritten with new information.

# Mapping functions

- Mapping functions determine how **memory blocks** are **placed** in the **cache**.
- Three **mapping** functions:
  - **Direct** mapping
  - **Associative** mapping
  - **Set-associative** mapping.


# Direct Mapped Cache

- Addressing scheme in **Direct Mapped Cache**:
  - cache block address = memory block address *mod* cache size
  - if cache size =  $2^m$ , cache address = lower m bits of n-bit memory address  
(equivalently: Cache address = Memory Address MOD cache size)
  - remaining upper n-m bits kept as **TAG bits** at each cache block
  - also need a **valid bit** to recognize valid entry



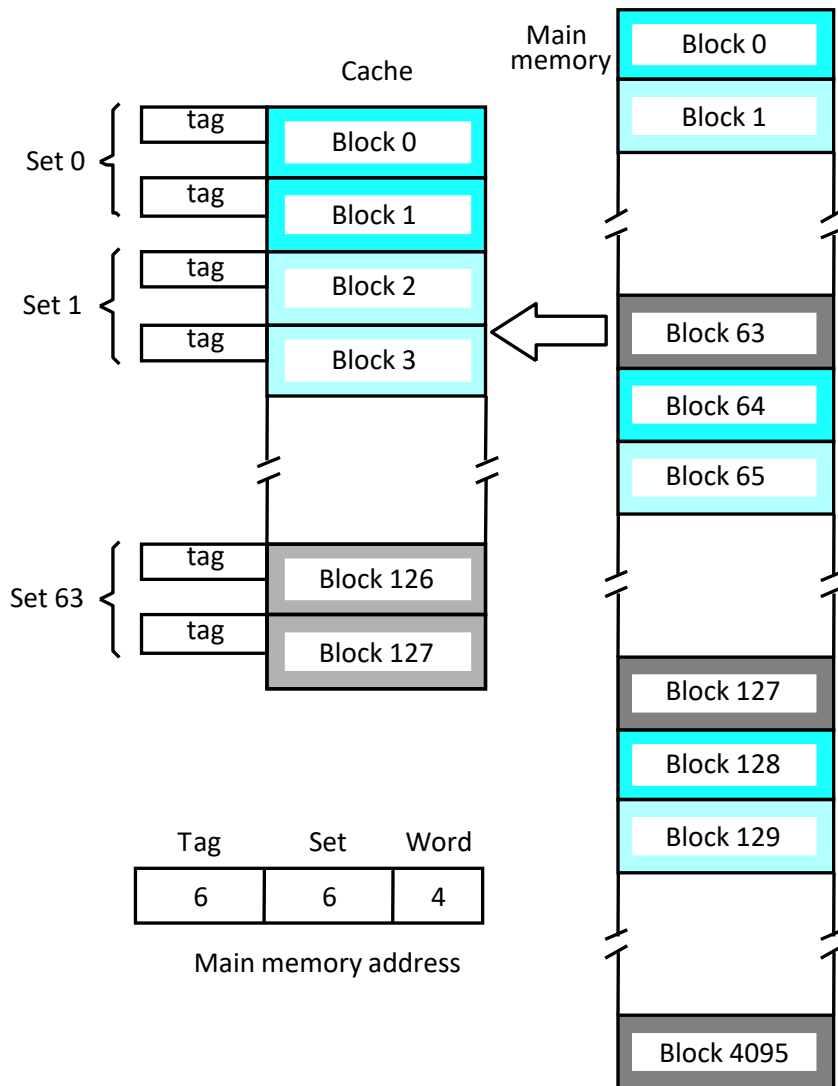
**Memory size = 32 ( $= 2^5$ ) blocks**  
**cache size = 8 ( $= 2^3$ ) blocks**  
So  $n=5$  and  $m=3$

# Associative mapping



- Main memory block can be placed into **any** cache position.
- Memory address is divided into **two** fields:
  - **(If Main memory is addressable by a 16-bit address)**
    - Low order **4 bits** identify the word within a block.
    - High order **12 bits** or **tag** bits identify a memory block when it is resident in the cache.
- Flexible, and uses cache space efficiently.
- Replacement algorithms can be used to replace an existing block in the cache when the cache is full.
- Cost is higher than direct-mapped cache because of the need to search all patterns to determine whether a given block is in the cache.

# Set-Associative mapping



Set-associative mapping combination of **direct** and **associative** mapping.

Blocks of cache are grouped into **sets**.

Mapping function allows a **block** of the main memory to **reside** in **any block** of a **specific set**. Assume that, No of **Blocks of memory** = 4096. Divide the cache into **64 sets**, with two blocks per set. (if cache size 128) Memory block **0, 64, 128** etc. map to **set 0**, and they can occupy either of the two positions.

16 bit Memory address is divided into **three** fields:

- Low order **4 bits** identify the **word** within a **block**.
- **6 bit** field determines the **set** number.
- High order **6 bit** fields are compared to the **tag** fields of the two blocks in a set.

Tag size = No of **Blocks of memory** / No of **sets of cache**  
 $= 4096 / 64 = 64 = 2^6$

Number of blocks per set is a design parameter.

- One **extreme** is to have **all** the **blocks** in **one set**, requiring no set bits (fully **associative** mapping).
- Other **extreme** is to have **one block** per set, is the same as **direct** mapping.

# Cache memories (contd..)

- A bit called as “**valid bit**” is provided for each **block**.
- If the block contains **valid** data, then the **bit** is set to **1**, else it is **0**.
- **Valid bits** are set to **0**, when the **power** is just turned **on**.
- When a **block** is loaded into the **cache** for the **first time**, the **valid** bit is set to **1**.
- **Data transfers** between main **memory** and **disk** occur directly **bypassing** the **cache**.
- When the **data** on a **disk** **changes**, the main **memory** **block** is also **updated**.
- However, if the **data** is also **resident** in the **cache**, then the **valid** bit is set to **0**.
- What happens if the **data** in the **disk** and main **memory** **changes** and the **write-back** protocol is being **used**?
- In this case, the **data** in the **cache** may also have **changed** and is indicated by the **dirty** bit.
- The **copies** of the **data** in the **cache**, and the main **memory** are **different**. This is called the **cache coherence problem**.
- One option is to **force** a **write-back** **before** the main **memory** is **updated** from the **disk**..

# Replacement algorithms

- **Direct-mapped** cache, the position that each memory block occupies in the cache is **fixed**.

As a result, the replacement strategy is **trivial**.

- **Associative** and **set-associative** mapping provide some flexibility in deciding which memory block occupies which cache block.

When a new block is to be transferred to the cache, and all the positions it may occupy are **full**, which block in the cache should be replaced?

- Locality of reference suggests that it may be okay to replace the block that has gone the longest time without being referenced.

This block is called Least Recently Used (**LRU**) block, and the replacement strategy is called **LRU replacement algorithm**.

LRU algorithm has been used extensively.

It provides poor performance in some cases.

Performance of the LRU algorithm may be improved by introducing a small amount of randomness into which block might be replaced.

Other replacement algorithms include removing the “**oldest**” block. However, they disregard the locality of reference principle and do not perform as well.